# Guided Practice on MongoDB Atlas

In this document we will work on MongoDB Atlas, a hosted MongoDB service to create a database, load documents and execute queries.

In a second part we will see how to work on larger databases with `mongosh` in the terminal.

At the end of this session, you should have a good understanding of MongoDB's query language (MQL).

Objectives:

- You connect to a MongoDB Atlas server
- You create a database and insert documents
- You build relatively complex queries on a MongoDB database

---

## Several ways to work with MongoDB

On your local machine

- MongoDB Community Edition Self-hosted, runs on your computer
- MongoDB Compass - the GUI tool for visualization
- `mongosh` the CLI (command line interface) in the terminal.

The MongoDB Shell, mongosh, is a JavaScript and Node.js **REPL** (Read-Eval-Print Loop) environment for interacting with MongoDB servers.

You can also work on MongoDB with your preferred scripting language: python, go, ruby, PHP, Java, etc...

Today, we are working on the hosted Atlas version.

- MongoDB Atlas provides a Cloud hosted service with free or paid offerings. **You don't have to install anything**.

MongoDB Atlas gives a hosted cluster on which we can create a database, import sample **collections** and understand how to perform CRUD operations in MongoDB.

> **CRUD** means Create, Read, Update and Delete: the four basic operations for a DBMS

---

## Atlas

Let's start by creating an account on Atlas https://www.mongodb.com/cloud/atlas/register

Then:

- Create a cluster and a project (I called my project `Sandbox`)

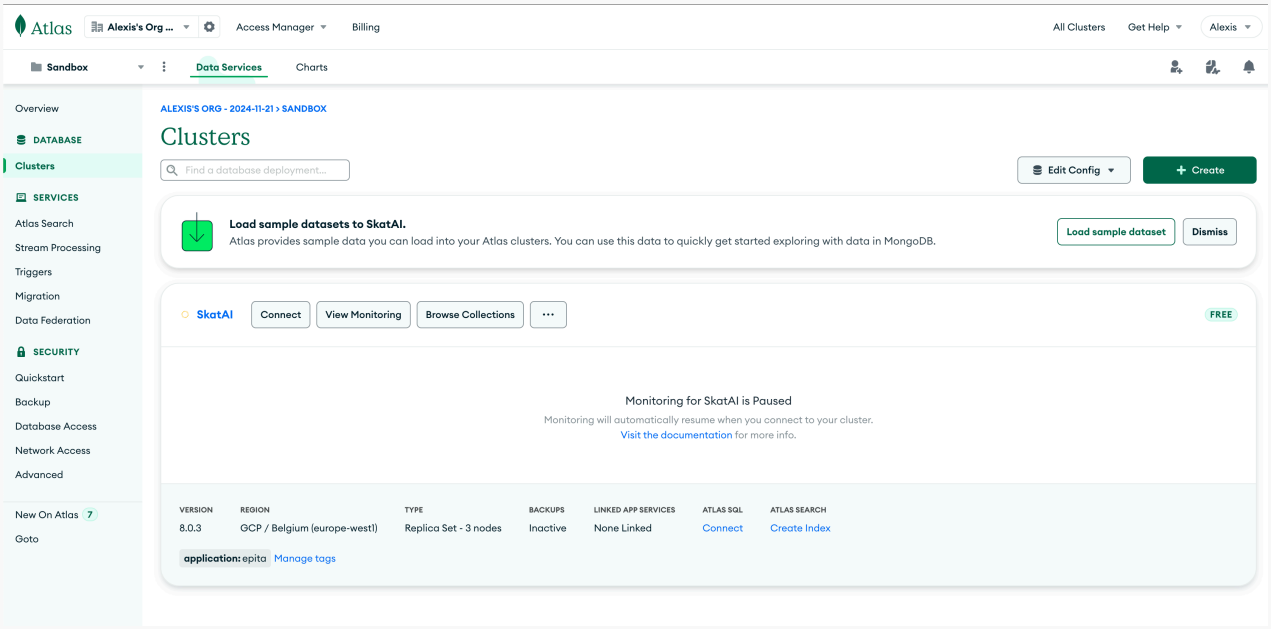


## Important

- **Copy your database password**.
- Check that your **IP** address is authorized. Visit `Network Access` (left navigation) and add it to the list
- To allow access from any IP add 0.0.0.0/0. This is not recommended for production but easier for our context.

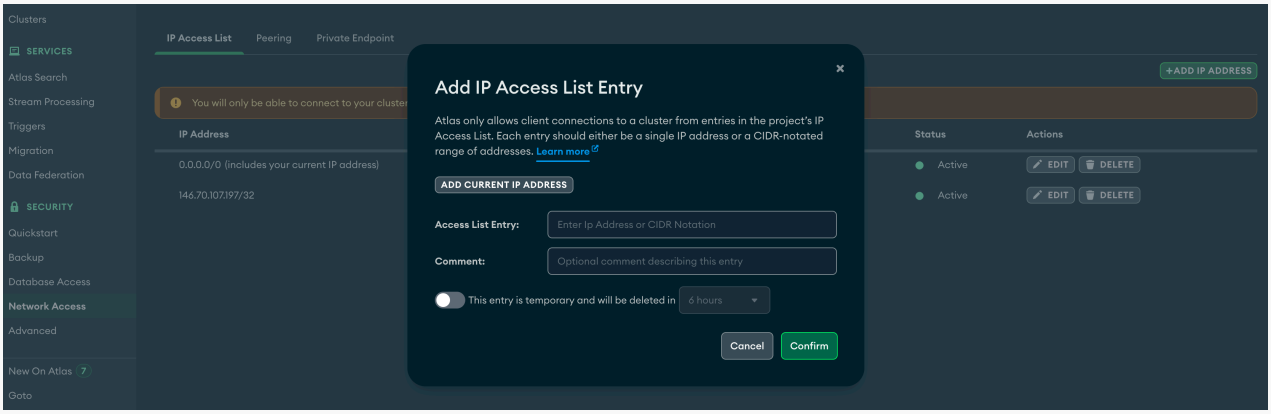


## Then create a MongoDB database

- Go to `clusters`
- Create a database: add the name `songsdb` and the collection name `songs`,

Note: do not confuse

- the project: `Sandbox`

- the database: `songsdb`
- the collection (table): `songs`

Then insert the following Document

```javascript
{
  "_id":{
    "$oid":"6745ab6f0e0bbdab062667c7"
    },
    "title": "Happy",
    "artist": "Pharrell Williams",
    "year": 2013,
    "mood": "joyful"
}
```

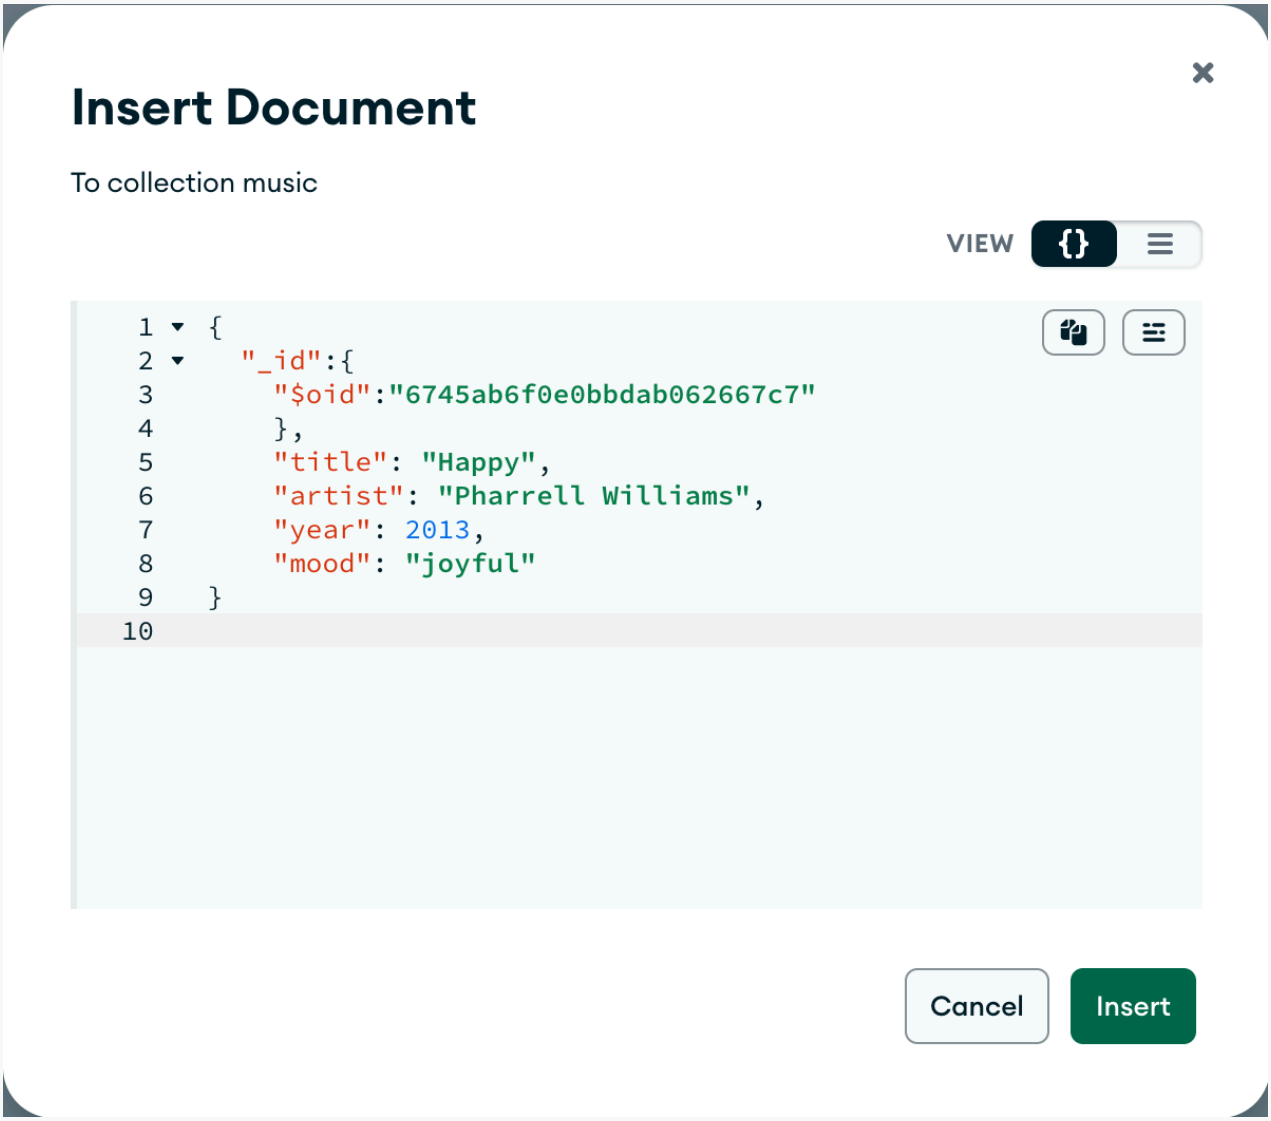


And insert another one

```JavaScript
{
  "_id":{
    "$oid":"abcdeab6f0e0bbdab062667aa"
  },
  "title": "Highway to hell",
  "artist": "AC/DC",
  "year": 1981,
  "mood": "energetic"
}
```

In both cases, the value of the `_id` is arbitrary. The tool proposes a default one with another value.

We therefore have 2 documents!

2 documents, wow that's mega exciting 🤣🤣🤣 !

Let's explore this MongoDB database.

---

## Queries in MongoDB

In MongoDB, writing a query comes down to writing `JSON`

| json | query |
|---|---|
| `{}` | returns all documents |
| `{ field : value }` | where field = value |
| `{ field : { $lt : value } }` | where field <= value (lt : less than) |

So if we want to find all documents in our `songs` collection, just write `{}` in the query field

This returns the 2 songs, the 2 documents.

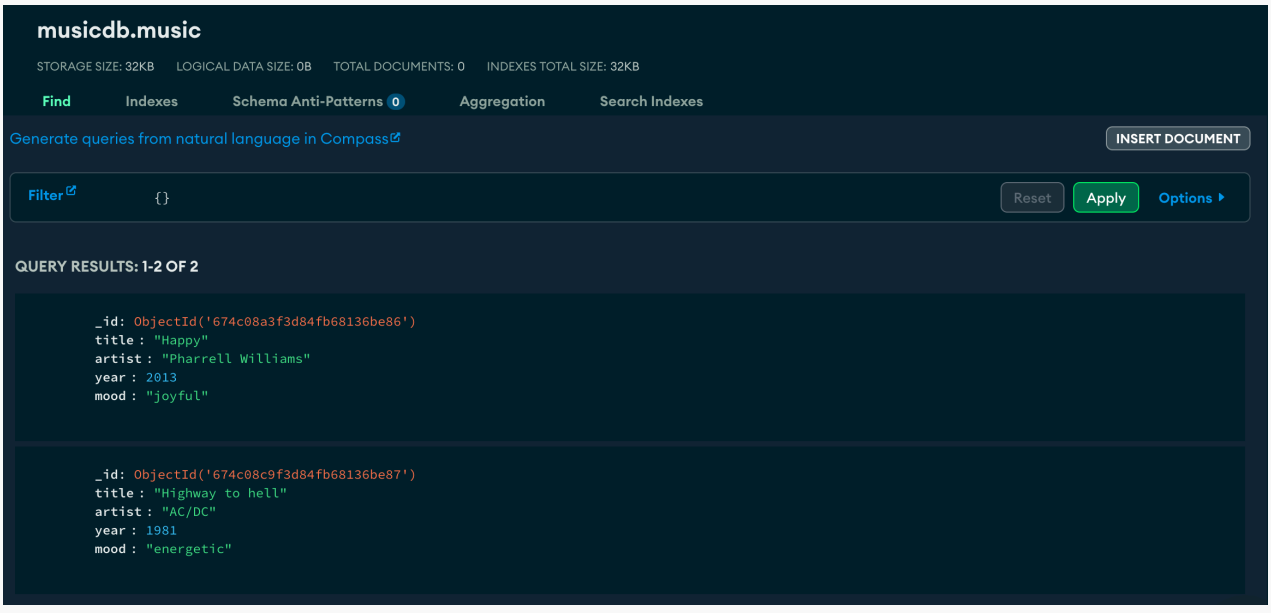

And the query `{ year: { $lt: 2000 } }` returns the song that was published before 2000.

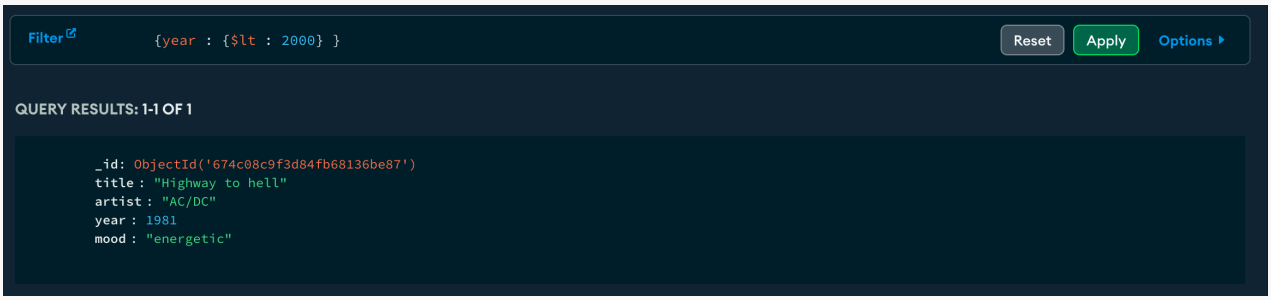


---

## Mooooooore data! (moaaaaaarrrr data)

2 documents is good but we need more data to really play.

Let's now import the sample datasets provided by ATLAS.

Go to `Clusters > dots > load sample dataset`

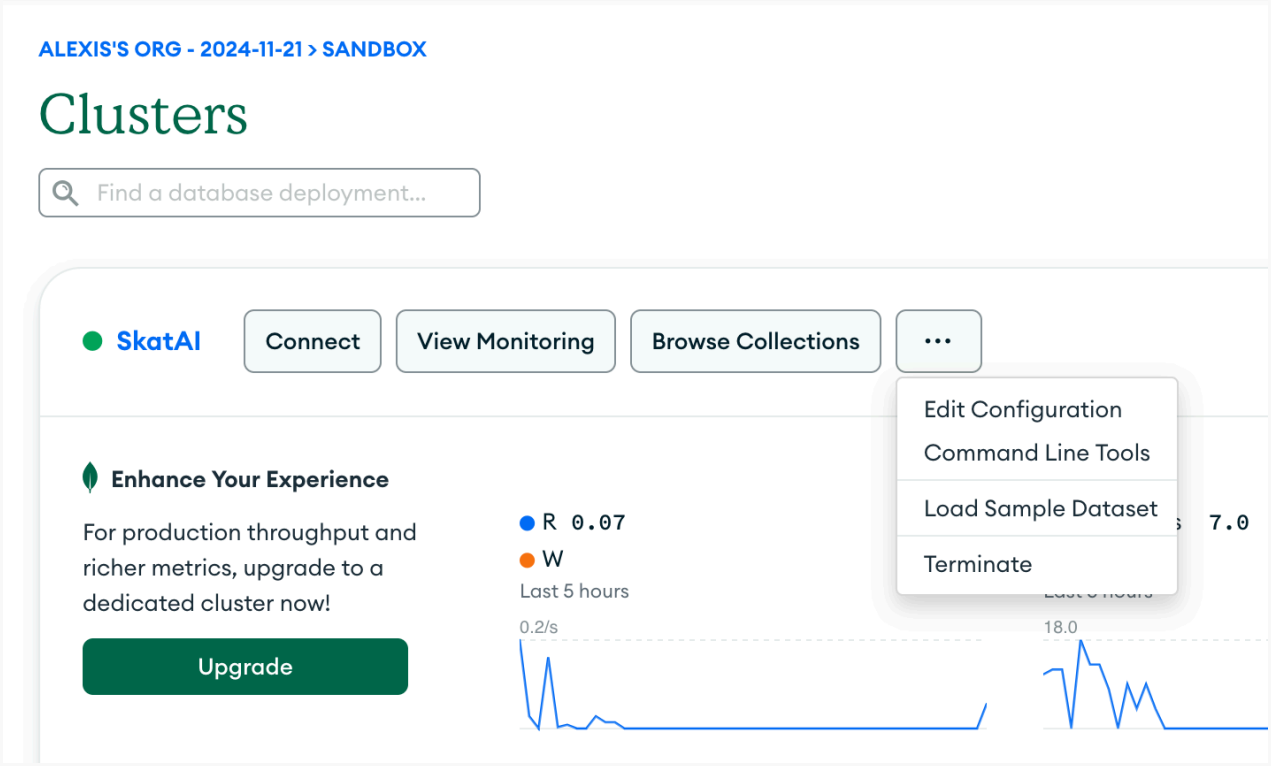


Then click on your project name (for me `Sandbox`) and `Browse collections` to see the available databases.

You get a new database `movies_mflix` with 5 collections.

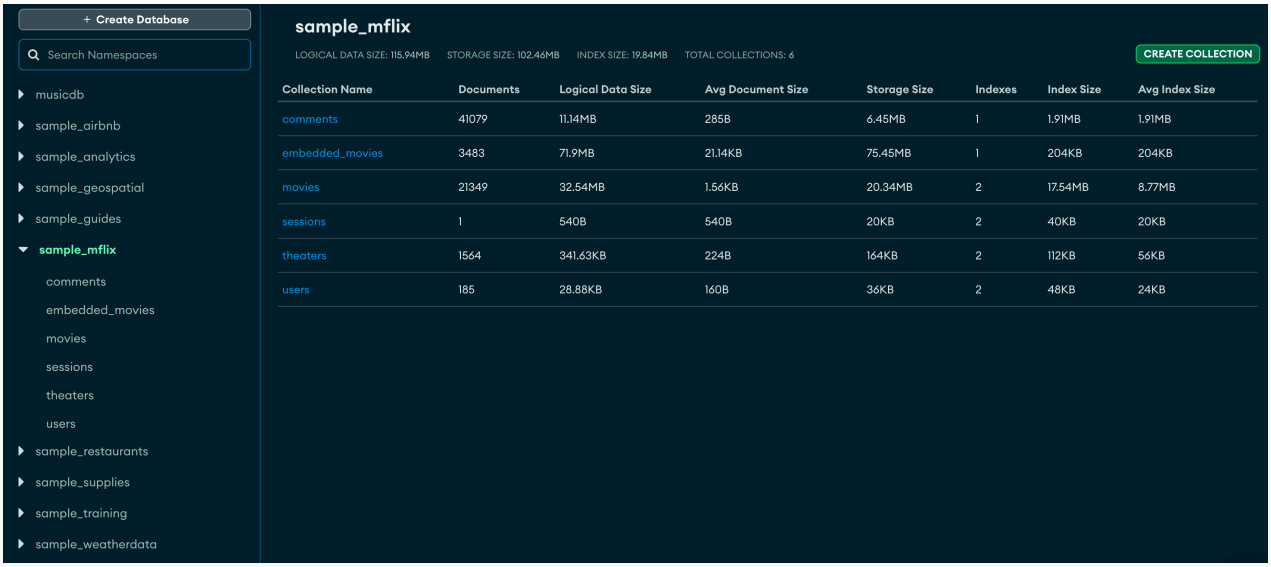


Look at the `embedded_movies` collection which contains 1525 documents and observe the structure of a document

**A document is a JSON record**

It can have:

- **nested JSON**: look at the nested dictionaries `imdb` and `tomatoes` fields.
- **arrays**: look at genres, cast, languages, writers, ...

**The primary key of a collection is always "`_id`".**

JavaScript

```
{
    "_id": {
        "$oid": "573a1390f29313caabcd5293"
    },
    "plot": "Young Pauline is left a lot of money when her wealthy uncle dies. H
    "genres": [
        "Action"
    ],
    "runtime": {
        "$numberInt": "199"
    },
    "cast": [
        "Pearl White",
        "Crane Wilbur",
        "Paul Panzer",
        "Edward Josè"
    ],
    "num_mflix_comments": {
        "$numberInt": "0"
    },
    "poster": "https://m.media-amazon.com/images/M/MV5BMzgxODk1Mzk2Ml5BMl5BanBnX
    "title": "The Perils of Pauline",
    "fullplot": "Young Pauline is left a lot of money when her wealthy uncle die
```

```
    "languages": [
        "English"
    ],
    "released": {
        "$date": {
            "$numberLong": "-1760227200000"
        }
    },
    "directors": [
        "Louis J. Gasnier",
        "Donald MacKenzie"
    ],
    "writers": [
        "Charles W. Goddard (screenplay)",
        "Basil Dickey (screenplay)",
        "Charles W. Goddard (novel)",
        "George B. Seitz",
        "Bertram Millhauser"
    ],
    "awards": {
        "wins": {
            "$numberInt": "1"
        },
        "nominations": {
            "$numberInt": "0"
        },
        "text": "1 win."
    },
    "lastupdated": "2015-09-12 00:01:18.647000000",
    "year": {
        "$numberInt": "1914"
    },
    "imdb": {
        "rating": {
            "$numberDouble": "7.6"
        },
        "votes": {
            "$numberInt": "744"
        },
        "id": {
            "$numberInt": "4465"
        }
    },
    "countries": [
        "USA"
    ],
    "type": "movie",
    "tomatoes": {
        "viewer": {
            "rating": {
                "$numberDouble": "2.8"
```

```
            },
            "numReviews": {
                "$numberInt": "9"
            }
        },
        "production": "Pathè Frères",
        "lastUpdated": {
            "$date": {
                "$numberLong": "1441993579000"
            }
        }
    }
}
```


The ATLAS user interface is well done but I prefer working with code rather than a web interface.

Let's move to python and `mongosh` to explore this movie database and learn how to make queries in MongoDB.

## Connecting via your language

Atlas allows you to connect to your cluster

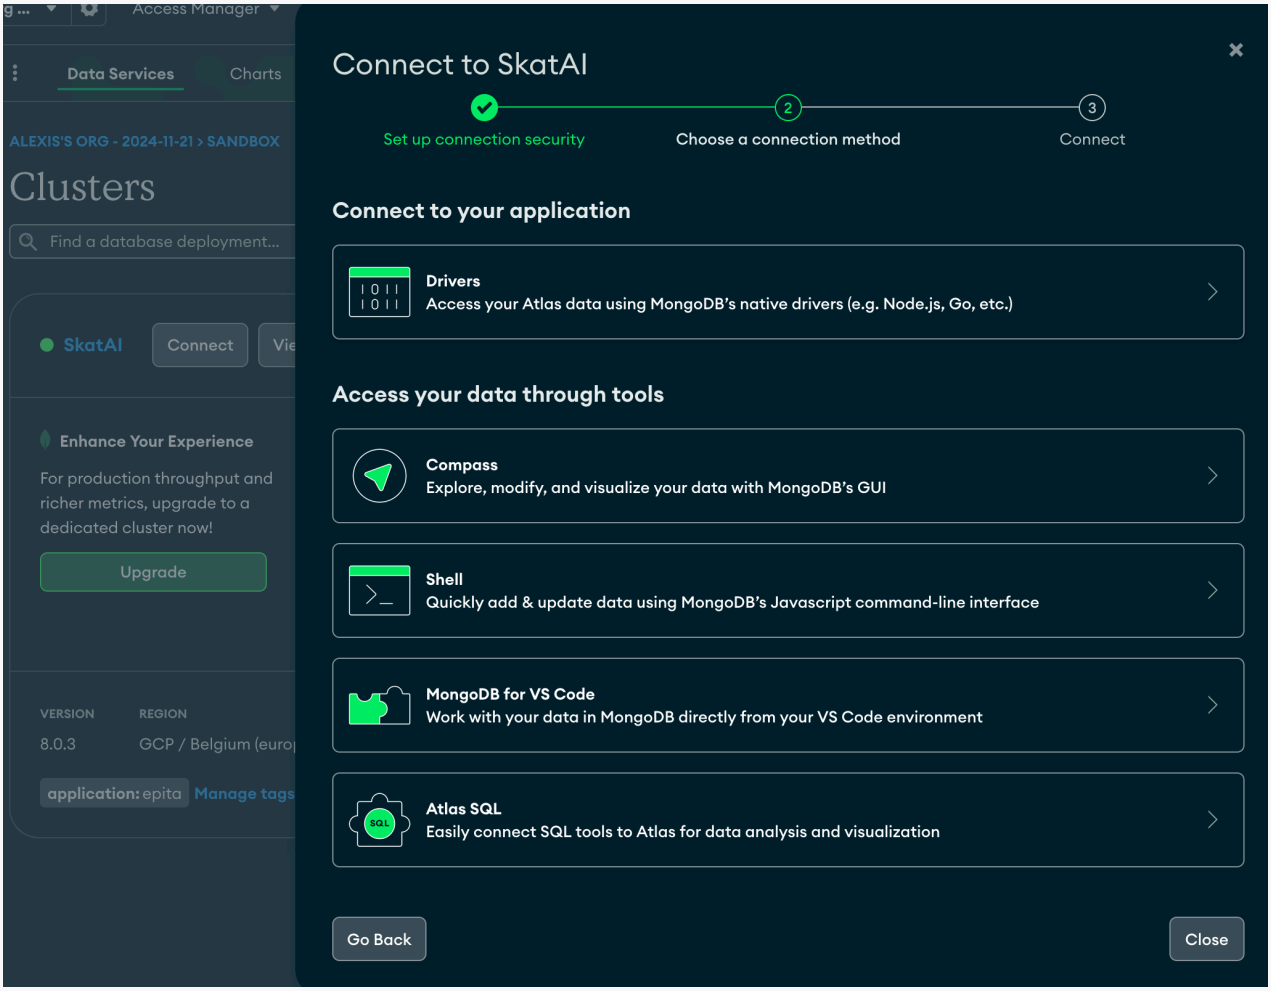

## Reset and password management

If you forgot your cluster password

## Connecting with Mongosh

Save the connection string in an environment variable

```
mongodb+srv://username:password@cluster_id.mongodb.net/
```

For example

```Bash
export MONGO_ATLAS_URI=mongodb+srv://alexis:password@skatai.w932a.mongodb.net/
```

Then connect with

```Bash
mongosh ${MONGO_ATLAS_URI}
```

### Connecting in Python

We need the pymongo package

```Bash
pip install pymongo
```

The connection string is

```Python
connection_string = "mongodb+srv://alexis:<db_password>@skatai.w932a.mongodb.net
```

Note: If you work in python, it's better to put the connection string as an environment variable (`$MONGO_ATLAS_URI`), in a `.env` file. then load it with `dotenv`

```Python
import os
from dotenv import load_dotenv

load_dotenv()

MONGO_ATLAS_URI = os.getenv('MONGO_ATLAS_URI')
```

And then instantiate the client with

```python
from pymongo import MongoClient

client = MongoClient(MONGO_ATLAS_URI)
```

Once we have a client, we can connect to the database

```python
db = client["sample_mflix"]
```

then instantiate a collection

```python
collection = db["movies"]
```

The collection is of class `pymongo.synchronous.collection.Collection` and has many methods:

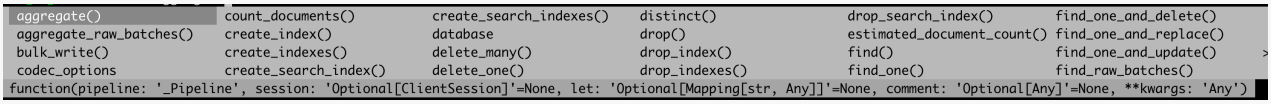


## Using the terminal with `mongosh`

How to install `mongosh`: https://www.mongodb.com/docs/mongodb-shell/install/

The language used in the MongoDB shell (`mongosh`) is **JavaScript**.

`mongosh` is an interactive JavaScript interface for MongoDB, allowing you to interact with your MongoDB instances, execute queries or perform administrative tasks in JavaScript.

`mongosh` also supports many JavaScript features including the use of variables, loops and functions.

Here are some examples of using JavaScript in `mongosh`:

- **Connect to a MongoDB instance:**

Local

```bash
mongosh "mongodb://localhost:27017"
```

On Atlas

```bash
mongosh mongodb+srv://alexis:password@skatai-ipsa.w932a.mongodb.net/
```

or

```bash
mongosh ${MONGO_ATLAS_URI}
```

Once connected to an instance

- **see the databases:**

```javascript
show dbs
```

- **Change database:**

```javascript
use myDatabase
```

- **see the collections:**

```javascript
show collections
```

- **Insert a document**

```javascript
db.songs.insertOne(
    {
        "title": "Green Onions",
      "artist": "Booket T and the MGs",
      "year": 1967,
      "mood": "[blues, soul]"
  }
)
```

```javascript
db.users.updateOne(
{ "artist": "Booket T and the MGs" },
{ $set: { "artist": "Booker T and the MGs" } }
);
```

```javascript

db.songs.insertOne(
    {
    "title": "Happy",
    "artist": "Pharrell Williams",
    "year": 2013,
    "mood": "joyful"
})
```

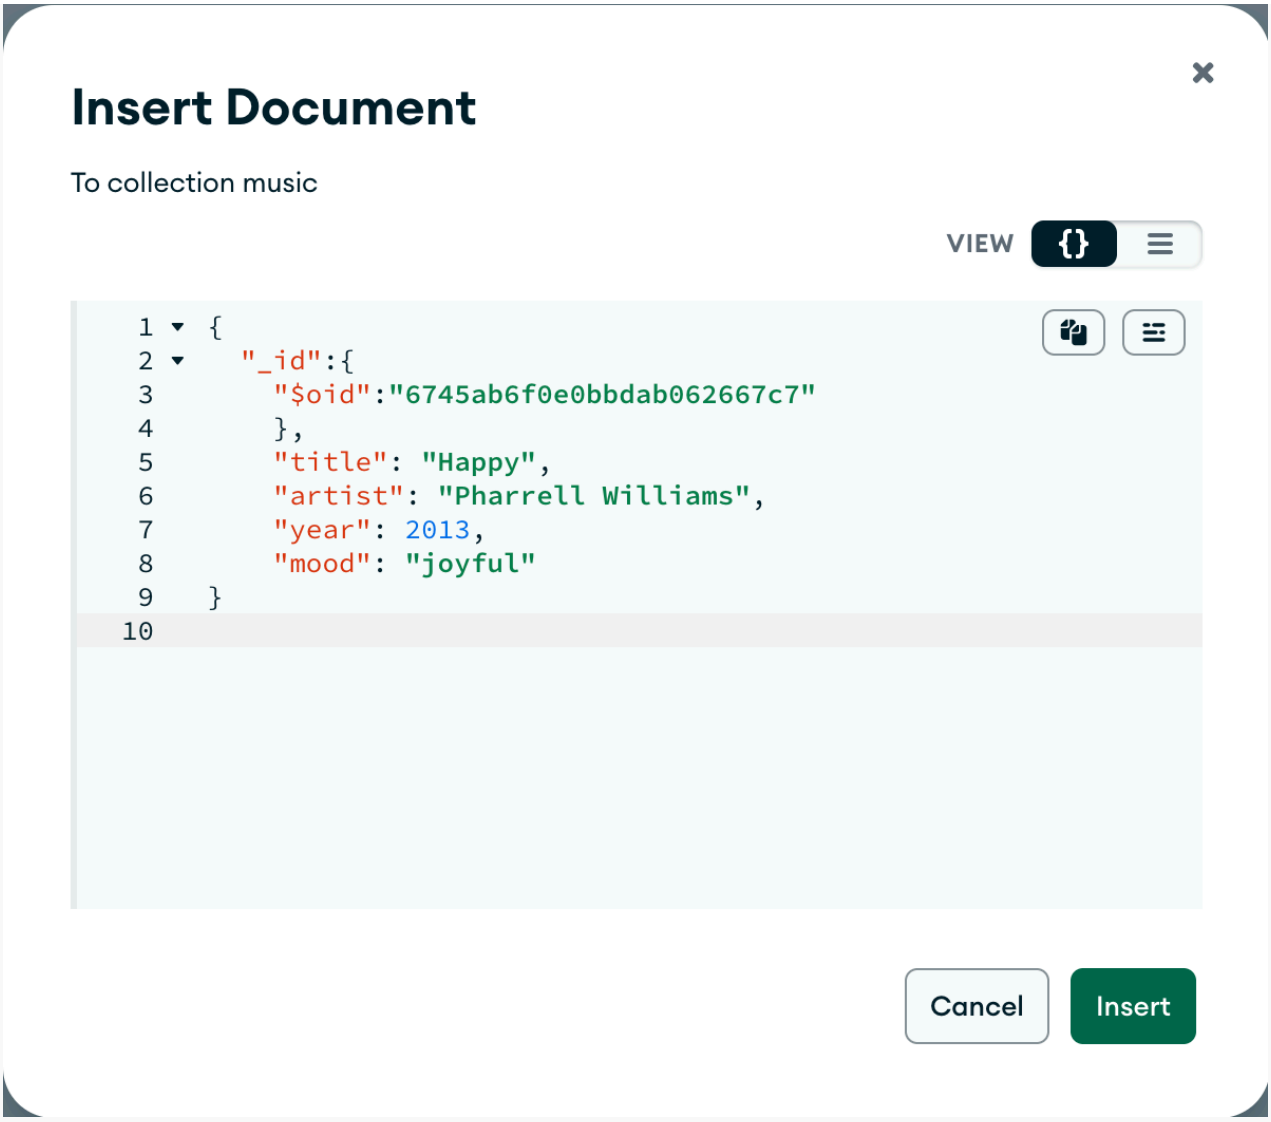


And insert another one

```JavaScript
db.songs.insertOne(
    {
  "title": "Highway to hell",
  "artist": "AC/DC",
  "year": 1981,
  "mood": "energetic"
})
```

- **Execute a query**

```JavaScript
db.songs.find({ year: { $lt: 2000 } })
```

- **Update a document**

```JavaScript
db.myCollection.updateOne({ name: "Alice" }, { $set: { age: 31 } })
```

- **Delete a document:**

```JavaScript
db.myCollection.deleteOne({ name: "Alice" })
```

- **Use JavaScript variables and functions**

```JavaScript
var name = "Alice";
db.myCollection.find({ name: name }).forEach(printjson);
```

`mongosh` is a powerful tool for querying and managing your MongoDB databases.

### Difference between queries in python and javascript

python: double quotes around fields and operators

```Python
db.movies.find(
    {"runtime": {"$gt" : 180}},  // Filter on movie duration
    { "_id": 0, "title": 1, "runtime": 1, "imdb.rating": 1 }  // Projection to i
)
```

`mongosh`: no need for quotes

```JavaScript
db.movies.find(
    {runtime: {$gt : 180}},  // Filter on movie duration
    { _id: 0, title: 1, runtime: 1, "imdb.rating": 1 }  // Projection to include
)
```

### Filtering

https://www.mongodb.com/docs/manual/reference/glossary/

The JSON that specifies the filtering arguments is called a **query predicate**. It's an expression that returns a boolean indicating whether a document matches the specified query.

For example, `{ title: { $eq: "Top Gun" } }`, which returns documents that have a "`title`" field whose value is "`Top Gun`".

An empty query predicate (`{ }`) returns all documents in the collection.

Main functions on collections

| function | returns |
| --- | --- |
| `find()` | all documents |
| `find_one()` | the 1st document |
| `distinct("<field>")` | list of distinct values for the `<field>` |
| `count_documents({})` | number of documents for the collection or returned by the filter in the query predicate |

Also note

- `find_one_and_replace()`
- `find_one_and_update()` and
- `delete_many()`
- `delete_one()`
- `drop_index()`
- etc ...

Note: You can also query the collection directly from the client with `db.<collection_name>.find()` without having to instantiate a `collection` object.

```Bash
collection.find({})
# or
db.movies.find({})
```

You can chain these methods with `limit` and `sort`

```Bash
db.movies.find({runtime: {$gt: 120}}).limit(3)
```

## Cursor

The result returned is a **cursor**.

```JavaScript
cursor = db.movies.find({})
```

A cursor is a pointer to a MongoDB query result set.

## Projection

In database language, **projecting** means selecting a subset of all possible fields.

In SQL, you simply list the column names

```SQL
select genres, plot from movies;
```

In MongoDB, you must specify the fields in a JSON object, right after the query predicate

```JavaScript
db.movies.find(
    {runtime: {$gt : 180}},  // Filter on movie duration
    { _id: 0, title: 1, runtime: 1, "imdb.rating": 1 }  // Projection to include
)
```

Here the projection is expressed by: `({ _id: 0, title: 1, runtime: 1, "imdb.rating": 1 })`

- `title: 1`: includes the title field.
- `runtime: 1,` includes the runtime
- `"imdb.rating": 1`: Includes the `imdb.rating` field.
- `_id: 0`: Excludes the `_id` field from the result (the default value is 1 if not specified).

The query returns

```JavaScript
  { runtime: 240, title: 'Napoleon', imdb: { rating: 7.4 } },
  { runtime: 281, title: 'Les Misèrables', imdb: { rating: 7.9 } },
  { runtime: 245, title: 'Flash Gordon', imdb: { rating: 7.3 } },
  { runtime: 238, title: 'Gone with the Wind', imdb: { rating: 8.2 } },
```

## Exercises

### In python

Let's execute some queries in Python or `mongosh` on the `movies` database

```Python
import os
from pymongo import MongoClient

connection_string = os.getenv('MONGO_ATLAS_URI')
client = MongoClient(connection_string)
db = client["sample_mflix"]
```

Then

- Retrieve the title and genres of movies that have the "Action" genre.

```python
cursor = db.movies.find(
    {"genres": "Action"},  # Filter: movies with 'Action' in the genres array
    {"_id": 0, "title": 1, "genres": 1}  # Projection: include title and genres,
)
```

to see the results with python

```python
for movie in cursor:
    print(movie)
```

**with** `mongosh`**:**

once connected

```javascript
cursor = db.movies.find(
    {"genres": "Action"},  # Filter: movies with 'Action' in the genres array
    {"_id": 0, "title": 1, "genres": 1}  # Projection: include title and genres,
)
```

### Number of documents

The fastest way to count documents is to use `count_documents`

```python
count = db.movies.count_documents({"imdb.rating": {"$gt": 8.0}})
```

> Note: in Python, you can clone the cursor to get its length and the number of returned documents. Cloning the cursor does not consume it

```python
len(list(cursor.clone()))
```

**with mongosh**

```javascript
db.movies.countDocuments({ "imdb.rating": { $gt: 8.0 } })
```

### Your turn

with

```python
cursor = db.movies.find( filter, projection).limit(5)
for movie in cursor:
    print(movie)
```

Write the filter and projection for the following queries and also return the number of documents with

`db.movies.count_documents(filter)`

- use projection to return only relevant fields or at minimum "title"
- limit results to 5 documents

1. Find movies with an IMDb rating greater than 8
   - filter: `{"imdb.rating": {"$gt" : 8}}`
   - projection: `{"_id": 0, "title": 1, "imdb.rating": 1}`
2. Movies released after 2000
3. Movies with a specific director: "Christopher Nolan". Show title, director and year
4. Retrieve movies with `tomatoes.viewer.rating > 4.0`, showing title and viewer rating.
5. Find movies that contain `"Comedy"` and `"Drama"` in the `genres` array. Use `{$all: [list of genres]}`
6. Combine a query with sorting: Retrieve the top 5 movies with the highest IMDb rating, showing title and rating. (you should only retrieve `imdb.rating` with `double` data type)
7. Movie query on a year range: Retrieve movies released between 1990 and 2000, showing title and year.
8. Movie query with missing fields: Find movies where the `fullplot` field does not exist. Use `$exists`.
9. Find all distinct genres
   - use `db.movies.distinct("genres")`
10. Movies with at least 2 genres
    - use `{"genres": {$size: 2}}`
11. Action movies, after 1950 with imdb ratings > 8, sort by year desc, imdb rating desc
    - use `{ "year": {$gt: 1950}, "imdb.rating": {$gt: 8}, "genres": "Action" }`
12. Movies with both genres: Action and Drama
    - `$and: [ {"genres": "Action"}, {"genres": "Drama"} ]`
13. Movies with either Action or Drama
    - `$or: [ {"genres": "Action"}, {"genres": "Drama"} ]`
14. Movies after 1950 with either imdb.rating > 0 or awards.wins > 5
    - use: `{ "year": {$gt: 1950}, $or: [ {"imdb.rating": {$gt: 8}}, {"awards.wins": {$gt: 5}} ] }`

---

## Conclusion

In this session you practiced:

- Setting up MongoDB Atlas, a cloud-hosted database service, including cluster creation and security configuration

- Writing basic MongoDB queries using JSON format:
  - Basic syntax: `{}` for all documents, `{field: value}` for equality, `{field: {$lt: value}}` for comparisons
  - How to query nested fields and arrays in complex documents
- Connecting to MongoDB Atlas using Python and `pymongo`:
  - Using basic operations: `find()`, `find_one()`, `distinct()`, `count_documents()`
  - Implementing projections to select specific fields
- Working with sample datasets (particularly the movies database) to practice:
  - Filtering and sorting data
  - Working with nested fields
  - Using operators like `$gt`, `$lt`, `$all`, `$exists`
  - Writing combined queries with multiple conditions

In the next session, we will dive deeper into MongoDB and look at more complex ways to query data using **aggregation pipelines**. We will also cover schema design and validation.

---

## To go further

For next time, you can:

- explore more Atlas sample databases and practice writing queries
- follow https://www.mongodb.com/docs/languages/python/pymongo-driver/current/read/ for more practice
- There are many free courses and tutorials in MongoDB University
  - Intro to MongoDB
  - CRUD in python

and much more